

When AI Exposure Rankings Disagree: Robust Top-25 Selection Across 661 U.S. Occupations

Ghassan Malkawi¹, Ahmed Abdelaziz Elsayed², Asem Omari¹
 Azmi Alazzam¹, Said Badreddine¹, Shaima Alharthi¹

¹Faculty of Computer Information Science, Higher Colleges of Technology
 Al Ain, Abu Dhabi, United Arab Emirates
 {gmalkawi,aomari,aalazzam,sbadreddine}@hct.ac.ae
 H00543372@hct.ac.ae

²Department of Computer Engineering and Computational Sciences, Canadian University Dubai
 Dubai, United Arab Emirates
 ahmed.elsayed@hud.ac.ae

Abstract—AI-exposure measures do not necessarily identify the same occupations as priorities. This creates a practical problem when institutions must choose a fixed set of occupations for screening, reskilling assessment, or further workforce analysis. Using 661 U.S. occupations and 18 coherent AI-exposure specifications, we ask which occupations remain stable priorities and which fixed Top-25 list is least sensitive to admitted measurement design. Under coherent designs, 14 occupations appear in every Top-25 and 37 appear in at least one; the more conservative independent interval box gives only 4 guaranteed and 87 possible members. A relative-minimax list raises worst-scenario capture from 93.70% for the baseline composite to 96.52%, although its exact independent-box absolute regret is slightly larger (0.0195 versus 0.0181). We derive exact best/worst rank certificates, a lossless existence-based screening theorem, and an exact fixed-list box-regret corner formula. A nested-support sensitivity then shows where membership conclusions change as the box contracts toward a declared point core. The reported values $g^* = 0.574320$, 0.434651 , and 0.078756 are transition boundaries: because best rank uses a strict comparison and worst rank a non-strict comparison, post-contraction membership applies immediately below the relevant boundary rather than being asserted at equality. The resulting Top-25 is a screening and further-assessment list, not a displacement forecast, transition probability, or program-success prediction.

Index Terms—AI exposure, occupational prioritization, robust Top- k , uncertain ranking, minimax regret, workforce screening

I. RESULTS FIRST: WHAT CHANGES FOR WORKFORCE SCREENING?

Different admitted AI-exposure constructions do not produce one unique occupational priority ranking. At $k = 25$, only 14 of 661 occupations occur in every coherent Top-25 list, while 37 occur in at least one. The independent coordinatewise interval box is intentionally more conservative: it certifies only 4 guaranteed members and 87 possible members. Table I converts those counts into a direct screening interpretation.

TABLE I
 OCCUPATIONAL STABILITY UNDER COHERENT AI-EXPOSURE
 DISAGREEMENT AT $k = 25$.

Status	Interpretation	n
Stable	In every coherent Top-25	14
Design-sensitive	In some but not all coherent Top-25 lists	23
Outside family	In no coherent Top-25 list	624




Fig. 1. Membership envelopes across evaluated selection sizes. Coherent specifications move all occupational scores together; the independent box permits coordinatewise adversarial movement.

When one fixed list is required, the relative-minimax solution retains at least 96.52% of each scenario’s own optimal Top-25 priority mass. The baseline composite retains at least 93.70%. Thus the applied question is not “which ranking is true?” but “which priorities are stable across admitted designs, and what fixed list limits worst-case loss when designs disagree?” The mathematics below supplies exact certificates for that decision.

II. DECISION PROBLEM AND COHERENT UNCERTAINTY

Let $i \in \{1, \dots, n\}$ index occupations, $s_i \geq 0$ be a scale variable, and q_i^t an uncertain modifier under coherent specification t . Priority is

$$p_i^t = s_i q_i^t, \quad (1)$$

with $p^t \in \mathbb{R}_{\geq 0}^n$. The application fixes $n = 661$ U.S. occupations. Scale is May 2025 OEWS employment share; occupational descriptors use O*NET 30.3. Eighteen primitive scenarios cross three exposure channels—AI Occupational Exposure (AIOE), Microsoft user-goal, and Microsoft AI-action/nonphysical—with six constructions: equal-five plus five leave-one-descriptor-out variants. These are coherent designs, not independent perturbations of individual occupations.

For each occupation,

$$\ell_i = \min_t p_i^t, \quad u_i = \max_t p_i^t, \quad (2)$$

and

$$\mathcal{P}_{\square} = \prod_{i=1}^n [\ell_i, u_i], \quad \mathcal{P}_T = \{p^t : t = 1, \dots, T\} \subseteq \mathcal{P}_{\square}. \quad (3)$$

The declared point core $c_i = p_i^{\text{AIOE, equal-five}}$ is one primitive scenario, hence

$$\ell_i \leq c_i \leq u_i \quad \forall i. \quad (4)$$

No probability distribution is assigned to either uncertainty set.

III. EXACT INDEPENDENT-BOX RANK CERTIFICATES

Tie treatment is asymmetric by construction. Best rank is existential and resolves endpoint ties in occupation i 's favor; worst rank is pessimistic and counts ties against i :

$$r_i^{\min} = 1 + \sum_{j \neq i} \mathbf{1}\{\ell_j > u_i\}, \quad (5)$$

$$r_i^{\max} = 1 + \sum_{j \neq i} \mathbf{1}\{u_j \geq \ell_i\}. \quad (6)$$

The endpoint realizations are feasible in $\mathcal{P}_{\square}$, so the certificates are exact:

$$i \text{ guaranteed Top-}k \iff r_i^{\max} \leq k, \quad (7)$$

$$i \text{ possible Top-}k \iff r_i^{\min} \leq k. \quad (8)$$

Figure 1 shows how the box and coherent envelopes differ over $k \in \{10, 25, 50, 100\}$.

IV. COHERENT FIXED-LIST RELATIVE MINIMAX

Let $V_k(p)$ be the sum of the k largest components of p , and define $V_t^* = V_k(p^t)$. For binary selection $x_i \in \{0, 1\}$ with $\sum_i x_i = k$, coherent capture and relative regret are

$$C_t(x) = \frac{\sum_i p_i^t x_i}{V_t^*}, \quad R_t(x) = 1 - C_t(x). \quad (9)$$

The fixed-list problem is

$$\min_{x, \rho} \rho \quad (10)$$

$$\text{s.t.} \quad \sum_i p_i^t x_i \geq (1 - \rho) V_t^*, \quad \forall t, \quad (11)$$

$$\sum_i x_i = k, \quad x_i \in \{0, 1\}, \quad 0 \leq \rho \leq 1. \quad (12)$$

This is a selecting-items minmax-regret problem; ρ is a worst relative shortfall, not a probability.

A. Lossless screening

Let $G = \{i : r_i^{\max} \leq k\}$ and $Z = \{i : r_i^{\min} > k\}$.

Theorem 1 (existence-based lossless screening). Assume $\mathcal{P}_T \subseteq \mathcal{P}_{\square}$. For any minimization objective $F(m_1, \dots, m_T)$ that is coordinatewise nonincreasing in scenario-wise selected masses $m_t = \sum_i p_i^t x_i$, there exists an optimum containing every item in G and no item in Z .

Proof sketch. If a size- k list omits $g \in G$, then at most $k - 1$ competitors satisfy $u_j \geq \ell_g$, so the list contains some h with $u_h < \ell_g$. Replacing h by g weakly increases every scenario-wise selected mass. If selected $i \in Z$, at least k competitors satisfy $\ell_j > u_i$, so some omitted j strictly dominates i in every scenario. Replacing i by j strictly increases every scenario-wise selected mass; because the family of size- k lists is finite, cycling is impossible. Repeating the exchanges gives an optimum containing G and excluding Z . $\square$

At $k = 25$ the screen fixes 4 occupations in and 574 impossible occupations out without changing the coherent optimum.

V. EXACT INDEPENDENT-BOX REGRET OF A FIXED LIST

For a fixed set $X = \{i : x_i = 1\}$, $|X| = k$, define absolute box regret

$$R_{\square}^{\text{abs}}(X) = \max_{p \in \mathcal{P}_{\square}} \left[V_k(p) - \sum_{i \in X} p_i \right]. \quad (13)$$

The maximizing corner is

$$v_i(X) = \begin{cases} \ell_i, & i \in X, \\ u_i, & i \notin X, \end{cases} \quad (14)$$

so the exact closed form is

$$R_{\square}^{\text{abs}}(X) = V_k(v(X)) - \sum_{i \in X} \ell_i. \quad (15)$$

Writing the expression with V_k avoids any ambiguity over the identity of a tied Top- k corner set.

VI. NESTED-SUPPORT CONTRACTION AND EXACT BOUNDARIES

We use possibility theory only as optional parameterization language; the data identify support and a point core,

TABLE II
REPORTED TOP-25 STRUCTURAL CONTRACTION BOUNDARIES.

Event	g^*	$1 - g^*$	Interpretation
Guaranteed $\rightarrow$ 14	.574320	.425680	coherent intersection
Possible $\rightarrow$ 37	.434651	.565349	cardinality; 32 shared
Full 25/25	.078756	.921244	point list

TABLE III
FIXED-LIST PERFORMANCE AT $k = 25$ OVER 18 PRIMITIVE SCENARIOS.

Rule	Mean cap.	Min. cap.	Max rel. regret	Box abs. regret
Scale only	93.02%	88.35%	11.65%	.0238
Baseline composite	96.14%	93.70%	6.30%	.0181
Relative minimax	96.99%	96.52%	3.48%	.0195

not a calibrated possibility distribution. Retained support width $g \in [0, 1]$ defines

$$L_i(g) = c_i - g(c_i - \ell_i), \quad U_i(g) = c_i + g(u_i - c_i). \quad (16)$$

By (4), this is a nested family with $g = 1$ equal to the full box and $g = 0$ equal to the point core.

Rank vectors are piecewise constant. A best-rank comparison changes only at $L_j(g) = U_i(g)$,

$$g_{ji}^{(B)} = \frac{c_j - c_i}{(c_j - \ell_j) + (u_i - c_i)}, \quad (17)$$

and a worst-rank comparison only at $U_j(g) = L_i(g)$,

$$g_{ji}^{(W)} = \frac{c_i - c_j}{(u_j - c_j) + (c_i - \ell_i)}, \quad (18)$$

when the denominator is positive and the ratio lies in $[0, 1]$. Thus all events follow from a finite exact breakpoint set rather than Monte Carlo sampling.

The endpoint convention matters. Since best rank uses $L_j > U_i$ while worst rank uses $U_j \geq L_i$, g^* is reported as the *transition boundary*. When contraction proceeds from $g = 1$ toward zero, the post-contraction membership state applies for $g < g^*$; we do not assert that a newly strict best-rank comparison already holds at equality. Table II therefore reports boundaries rather than “first-at-equality” events.

For any strictly decreasing homeomorphism $\phi : [0, 1] \rightarrow [0, 1]$ with $\phi(0) = 1$ and $\phi(1) = 0$, setting $g = \phi(\alpha)$ traverses the same nested supports in the same order. Examples $\phi(\alpha) = 1 - \alpha$ and $\phi(\alpha) = (1 - \alpha)^2$ change numerical α labels but not the structural event sequence.

VII. APPLICATION RESULTS

Table III compares three fixed lists. The relative-minimax list improves the minimum coherent-scenario capture to 96.52%, from 93.70% for the baseline composite and 88.35% for scale only. Its exact box absolute regret is 0.0195, slightly larger than the baseline’s 0.0181, so no list is unconditionally “most robust” without naming the uncertainty structure and criterion.




Fig. 2. Three reported exact Top-25 contraction boundaries. Points mark structural transition values; the membership state after a strict boundary crossing is interpreted according to the tie convention in the text.




Fig. 3. Robustness trade-off at $k = 25$. Lower is better on both axes; the baseline and relative-minimax lists are non-dominating across uncertainty structures.

Criterion normalization also matters. Define coherent absolute regret

$$R_t^{\text{abs}}(x) = V_t^* - \sum_i p_i^t x_i, \quad (19)$$

and solve $\min_{x: \sum_i x_i = k} \max_t R_t^{\text{abs}}(x)$. This alternative changes two of the 25 occupations: its minimum relative capture is 94.79% and maximum relative regret 5.21%, with 23/25 overlap with the relative-minimax list.

VIII. PRACTICAL INTERPRETATION AND LIMITS

The results show why occupational AI prioritization should not be treated as one fixed ranking. Some occupations remain priority candidates across all 18 admitted constructions, while 23 additional occupations enter or leave the Top-25 depending on measurement design. For analysts or institutions that require one fixed screening list, the relative-minimax solution is a defensible compromise because it maximizes the worst retained share of scenario-optimal priority mass across the declared designs.

This interpretation is deliberately limited. The list is intended for screening and further assessment. It does not estimate AI-caused displacement, worker movement, training eligibility, transition capacity, or program success. Employment is a scale variable rather than a welfare weight. The 18-scenario family is finite and reflects only the admitted exposure/construction choices; the interval box is broader because it permits jointly unattainable coordinate combinations. A separate SIPP diagnostic in the Supplement is descriptive only.

IX. CONCLUSION

Across 661 U.S. occupations, AI-exposure disagreement materially affects who enters a Top-25 priority list. Yet the disagreement is structured: 14 occupations are stable across all coherent designs and 37 form the coherent candidate envelope. Exact interval ranks distinguish guaranteed from merely possible membership, while coherent minimax supplies one fixed list when a decision must be made. Relative minimax raises worst-scenario capture from 93.70% to 96.52% compared with the baseline composite, without claiming universal robustness. Exact breakpoint semantics further prevent equality/tie conventions from being mistaken for new membership at a boundary. The contribution is therefore an applied occupational-screening result supported by exact optimization certificates, rather than an optimization method presented with a workforce example.

REFERENCES

- [1] D. Bertsimas and M. Sim, "Robust discrete optimization and network flows," *Math. Programming*, vol. 98, pp. 49–71, 2003.
- [2] D. Bertsimas and M. Sim, "The price of robustness," *Operations Research*, vol. 52, no. 1, pp. 35–53, 2004.
- [3] D. Bertsimas, D. B. Brown, and C. Caramanis, "Theory and applications of robust optimization," *SIAM Review*, vol. 53, no. 3, pp. 464–501, 2011.
- [4] H. Aissi, C. Bazgan, and D. Vanderpooten, "Min–max and min–max regret versions of combinatorial optimization problems: A survey," *EJOR*, vol. 197, no. 2, pp. 427–438, 2009.
- [5] J. Ide and A. Schöbel, "Robustness for uncertain multi-objective optimization," *OR Spectrum*, vol. 38, pp. 235–271, 2016.
- [6] B. Roy, "Robustness in operational research and decision aiding," *EJOR*, vol. 200, no. 3, pp. 629–638, 2010.
- [7] M. A. Soliman and I. F. Ilyas, "Ranking with uncertain scores," in *Proc. IEEE ICDE*, 2009, pp. 317–328.
- [8] K. Mouratidis, J. Zhang, and H. Pang, "Maximum rank query," *PVLDB*, vol. 8, no. 12, pp. 1554–1565, 2015.
- [9] K. Mouratidis and B. Tang, "Exact processing of uncertain top-k queries in multi-criteria settings," *PVLDB*, vol. 11, no. 8, pp. 866–879, 2018.
- [10] OECD/JRC, *Handbook on Constructing Composite Indicators*. OECD Publishing, 2008.
- [11] E. W. Felten, M. Raj, and R. Seamans, "Occupational, industry, and geographic exposure to artificial intelligence," *Strategic Management Journal*, vol. 42, no. 12, pp. 2195–2217, 2021.
- [12] K. Tomlinson *et al.*, "Working with AI: Measuring the applicability of generative AI to occupations," arXiv:2507.07935, 2025.
- [13] M. Yin, H. Vu, and C. Persico, "How (un)stable are LLM occupational exposure scores?" NBER Working Paper 35110, 2026.
- [14] C. F. Manski, *Partial Identification of Probability Distributions*. Springer, 2003.
- [15] S. Greco, V. Mousseau, and R. Słowiński, "Ordinal regression revisited," *EJOR*, vol. 191, no. 2, pp. 416–436, 2008.
- [16] S. Steegen *et al.*, "Increasing transparency through a multiverse analysis," *Perspectives on Psychological Science*, vol. 11, no. 5, pp. 702–712, 2016.
- [17] U. Simonsohn, J. P. Simmons, and L. D. Nelson, "Specification curve analysis," *Nature Human Behaviour*, vol. 4, pp. 1208–1214, 2020.
- [18] I. Averbakh, "Computing and minimizing the relative regret in combinatorial optimization with interval data," *Discrete Optimization*, vol. 2, no. 4, pp. 273–287, 2005.
- [19] P. Kouvelis and G. Yu, *Robust Discrete Optimization and Its Applications*. Kluwer, 1997.
- [20] L. J. Savage, "The theory of statistical decision," *JASA*, vol. 46, no. 253, pp. 55–67, 1951.
- [21] E. Conde, "An improved algorithm for selecting p items with uncertain returns," *Math. Programming*, vol. 100, no. 2, pp. 345–353, 2004.
- [22] A. Kasperski, *Discrete Optimization with Interval Data: Minmax Regret and Fuzzy Approach*. Springer, 2008.
- [23] M. Goerigk and A. Schöbel, "Algorithm engineering in robust optimization," in *Algorithm Engineering*, Springer, 2016, pp. 245–279.
- [24] Q. Huangfu and J. A. J. Hall, "Parallelizing the dual revised simplex method," *Math. Programming Computation*, vol. 10, no. 1, pp. 119–142, 2018.
- [25] A. Kasperski and P. Zieliński, "Minmax regret approach and optimality evaluation with interval and fuzzy weights," *EJOR*, vol. 200, no. 3, pp. 680–687, 2010.
- [26] D. Dubois and H. Prade, "Practical methods for constructing possibility distributions," *Int. J. Intelligent Systems*, vol. 31, no. 3, pp. 215–239, 2016.
- [27] U.S. Bureau of Labor Statistics, *Occupational Employment and Wage Statistics, May 2025 National Estimates*, 2026.
- [28] U.S. Department of Labor, ETA, *O*NET 30.3 Database*, May 2026 release.
- [29] R. Lahdelma, J. Hokkanen, and P. Salminen, "SMAA—stochastic multiobjective acceptability analysis," *EJOR*, vol. 106, no. 1, pp. 137–143, 1998.
- [30] M. Goerigk and M. Khosravi, "Robust combinatorial optimization problems under budgeted interdiction uncertainty," *OR Spectrum*, vol. 47, pp. 255–285, 2025.